

A
Project Report
On

Predicting Life Expectancy using Machine Learning

By

Rohit Tiwari

Email : rtiwari432@gmail.com

Web Page(User Interface) Link

https://node-red-mvldi.eu-gb.mybluemix.net/ui/#!/0?socketid=5vG-jW5aL_NYFmVJAAA0

1. INTRODUCTION

Life expectancy is the average number of years a person in a population could expect to live after age x . It is the life table parameter most commonly used to compare the survival experience of populations. The age most often selected to make comparisons is 0.0 (i.e., birth), although, for many substantive and policy analyses, other ages such as 65+ and 85+ are more relevant and may be used (e.g., for determining person-years of Medicare and Social Security benefit entitlement). To calculate life expectancy at age x (e_x), age-specific mortality and population counts are needed to determine the age-specific mortality rates (i.e., the q_x) and survival probabilities (l_x) used in life table computations. Life expectancy is determined by multiplying the sequence of the probabilities of survival at each age to determine the proportion of a population expected to survive to age x . The number of persons expected to be alive in each single year of age category after age x is summed to determine the total number of years left to be lived after the index age (L_x). The total number of person-years to be lived after age x divided by the expected number of survivors to that age yields the life expectancy at age x .

1.1 Overview

In this project, we have to create a new model based on the data provided to evaluate the life expectancy.

The data offers a timeframe from 2015 to 2022. The output algorithms have been used to test if they can maintain their accuracy in predicting the life expectancy for data they haven't been trained. Following algorithms have been used:

Linear Regression

Ridge Regression

Lasso Regression

Elastic Net Regression

Linear Regression with Polynomic features

Decision Tree Regression

Random Forest Regression

World Health Organization (WHO) keeps track of the health status as well as many other related factors for all countries. The data-sets are made available to public for the purpose of health data analysis. The data-set related to life expectancy, health factors for 193 countries has been collected from the same WHO data repository website and its corresponding economic data was collected from United Nation website.

1.2 Purpose

The purpose of the project is to design a model for predicting Life Expectancy rate of a country given various features such as year, GDP, education, alcohol intake of people in the country, expenditure on healthcare system and some specific disease related deaths that happened in the country are given.

2. LITERATURE SURVEY

2.1 Existing Problem

To Predict the Life expectancy of a country based on various factor such as GDP, BMI, HIV/AIDS, Year, Alcohol intake and etc.

2.2 Proposed solution

We are using Machine Learning Algorithm called Random forest Regressor to solve this problem.

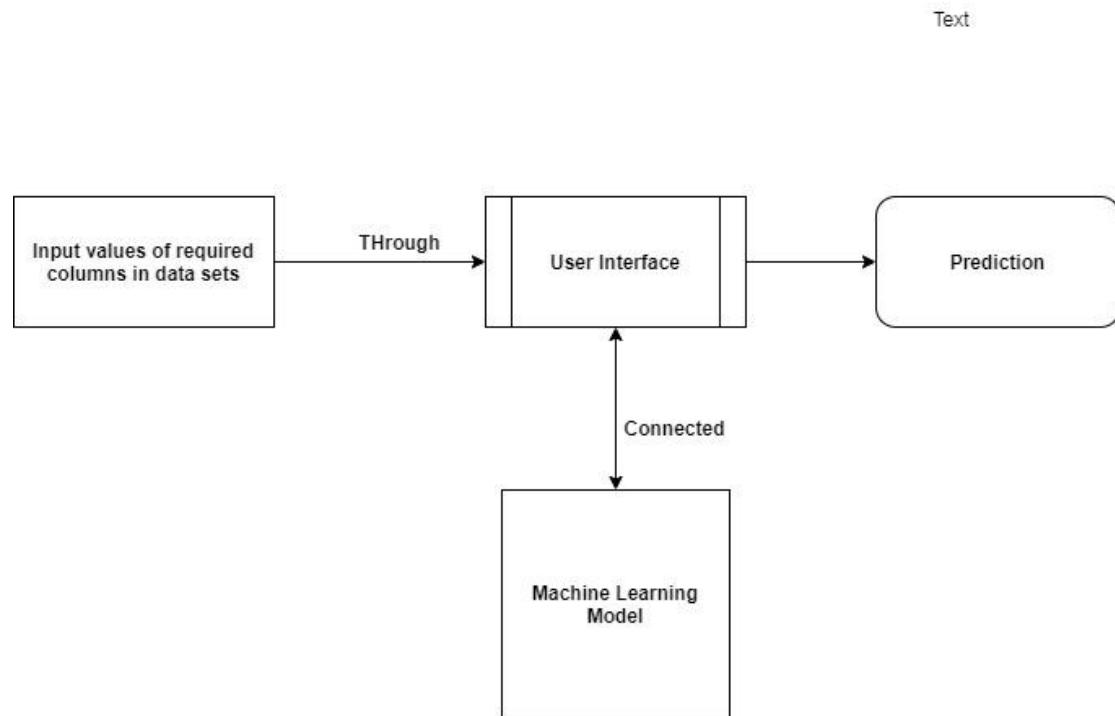
We are using data set provided by WHO.

Steps

- Import Data set
- Fill the values of columns if they are empty.
- Convert all String values in to integer or float.
- Drop the unnecessary columns.
- Import and train the model.
- Cheack accuracy of model.
- Predict the life Expectency of the Country.

3. THEORETICAL ANALYSIS

3.1 Block Diagram



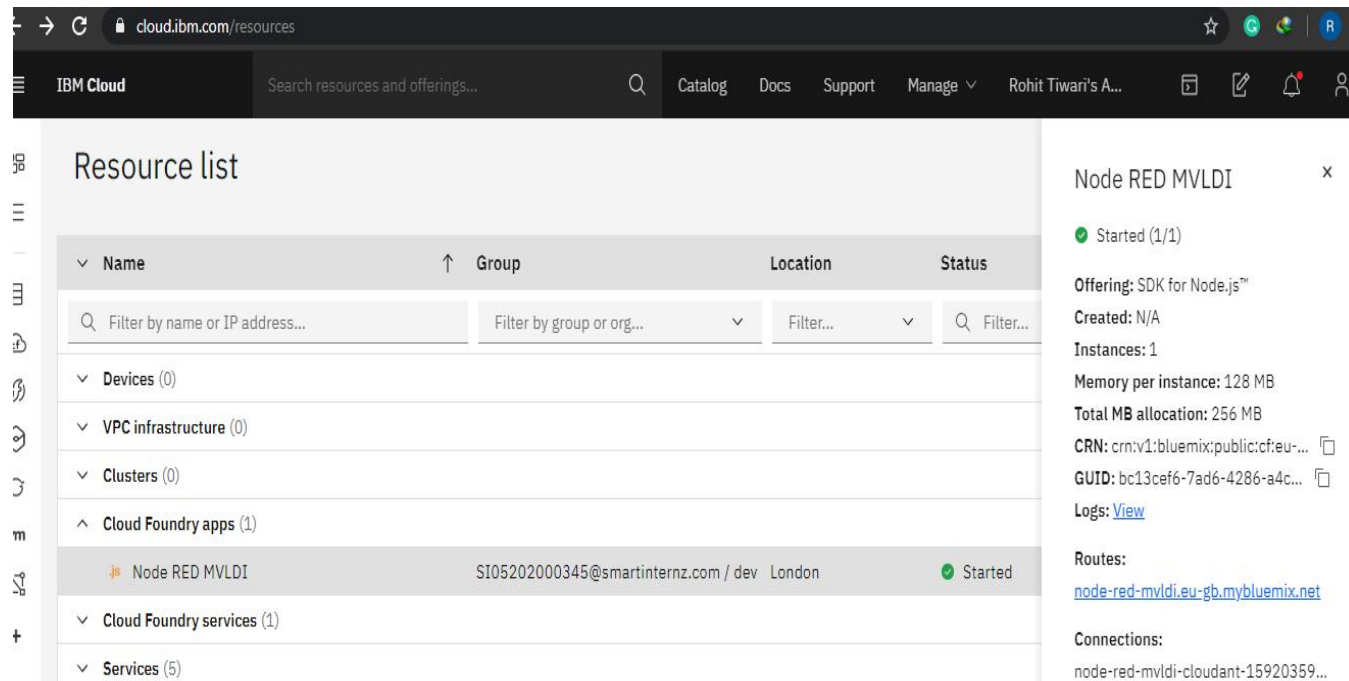
3.2 Hardware / Software designing

8 GB RAM

Python, IBM Cloud, IBM Watson, ML, WATSON Studio, Node-Red.

4. EXPERIMENTAL INVESTIGATIONS

A) IBM Cloud Resource List



The screenshot shows the IBM Cloud 'Resource list' page. The main table lists resources with columns for Name, Group, Location, and Status. The 'Node RED MVLDI' resource is highlighted. The sidebar on the right provides details for this resource.

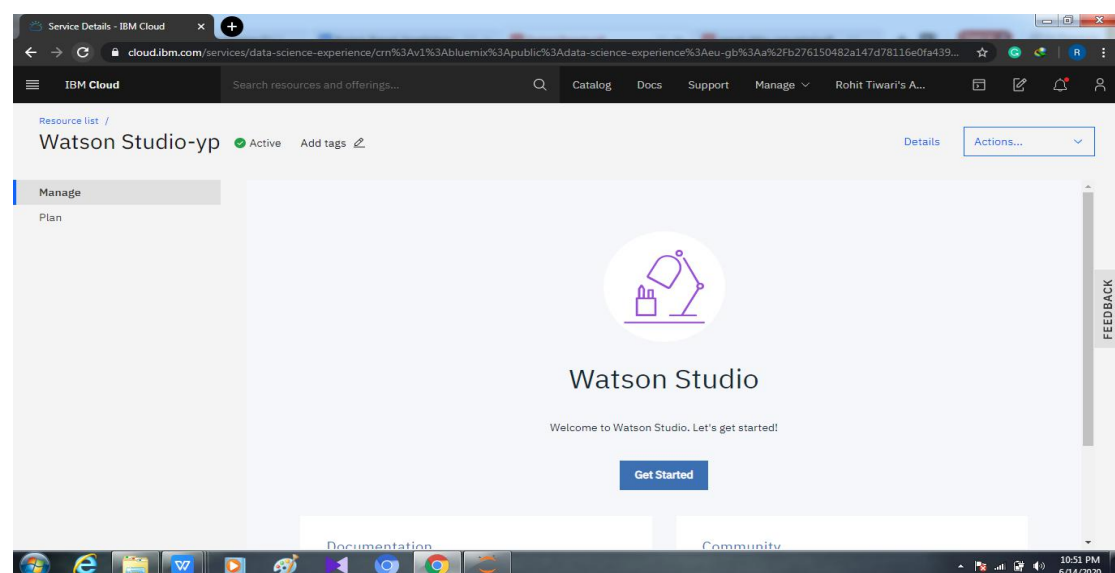
Name	Group	Location	Status
Node RED MVLDI	SI05202000345@smartinternz.com / dev	London	Started

Node RED MVLDI Details:

- Started (1/1)
- Offering: SDK for Node.js™
- Created: N/A
- Instances: 1
- Memory per instance: 128 MB
- Total MB allocation: 256 MB
- CRN: crn:v1:bluemix:public:cf:eu-...
- GUID: bc13cef6-7ad6-4286-a4c...
- Logs: [View](#)
- Routes: [node-red-mvldi.eu-gb.mybluemix.net](#)
- Connections: node-red-mvldi-cloudant-15920359...

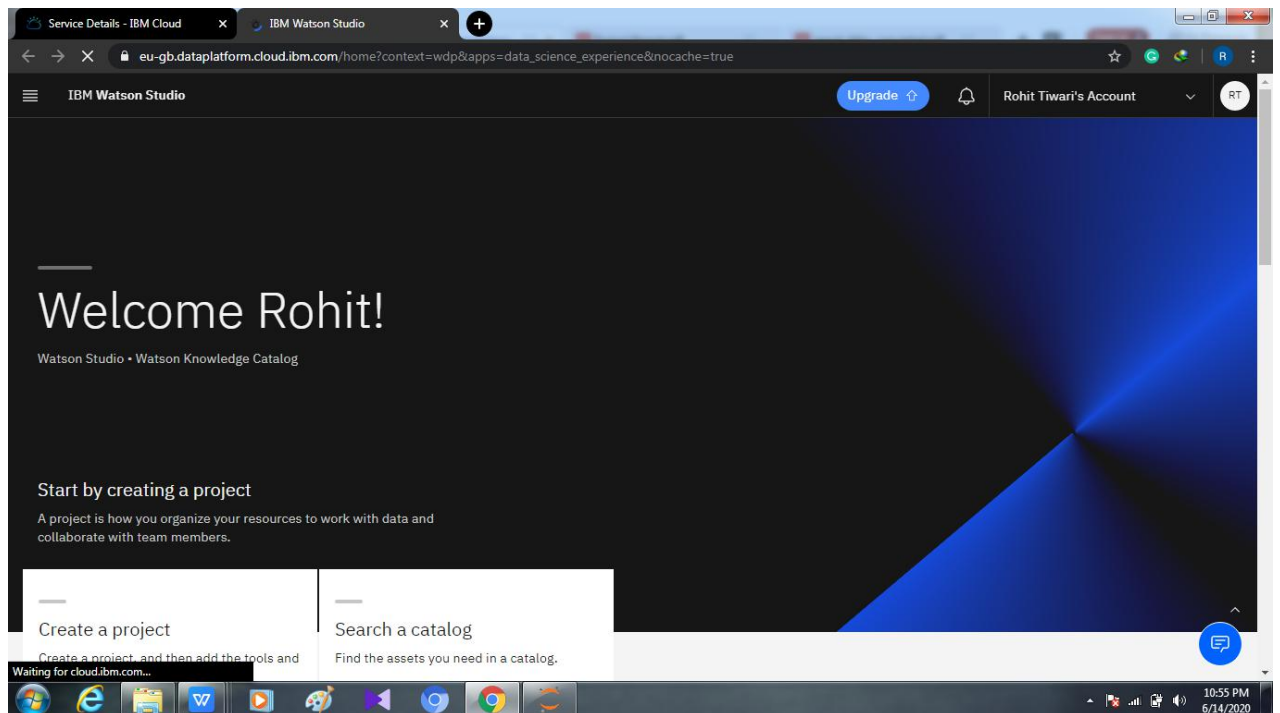
B) IBM Watson Studio

i)

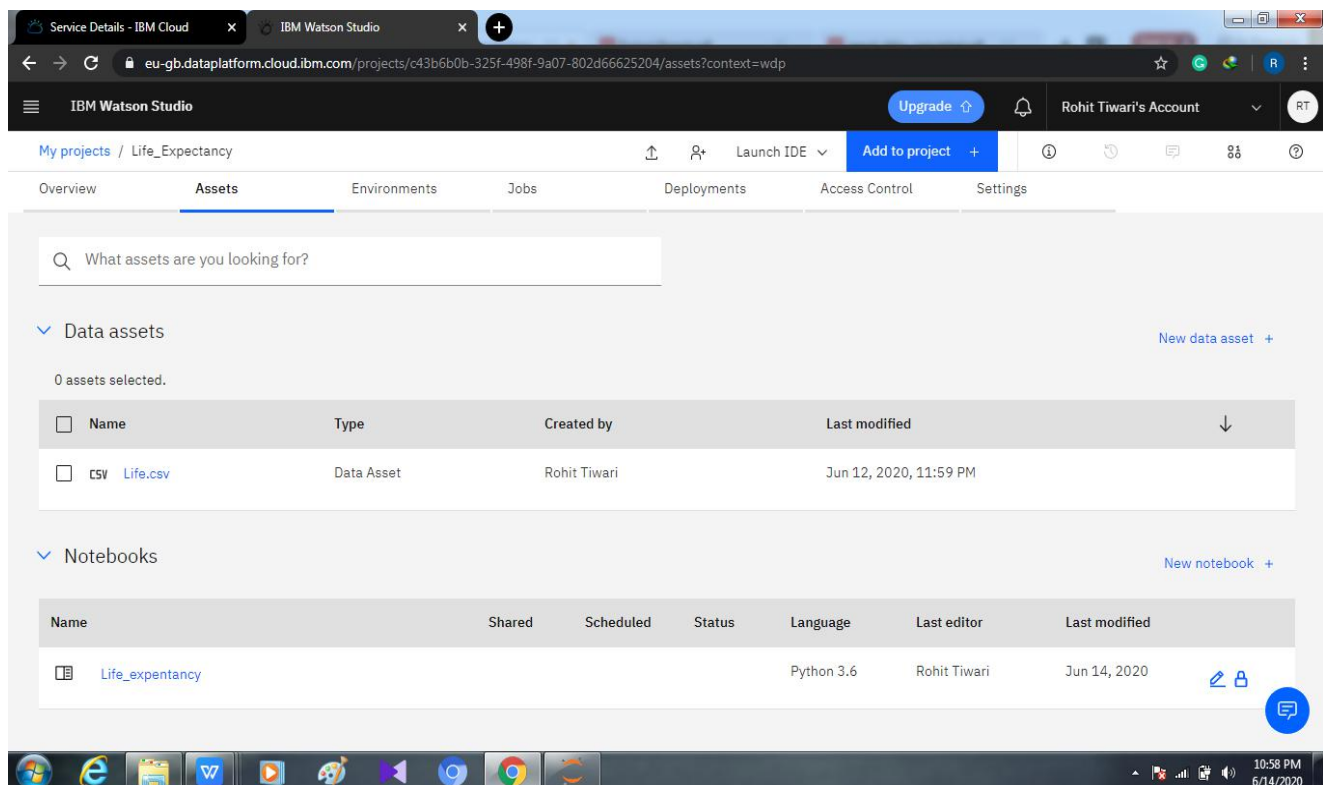


The screenshot shows the IBM Watson Studio 'Get Started' page. The page features the Watson Studio logo and a 'Get Started' button. The sidebar on the left shows 'Manage' and 'Plan' options. The top navigation bar includes 'Service Details - IBM Cloud' and 'cloud.ibm.com/services/data-science-experience/...'.

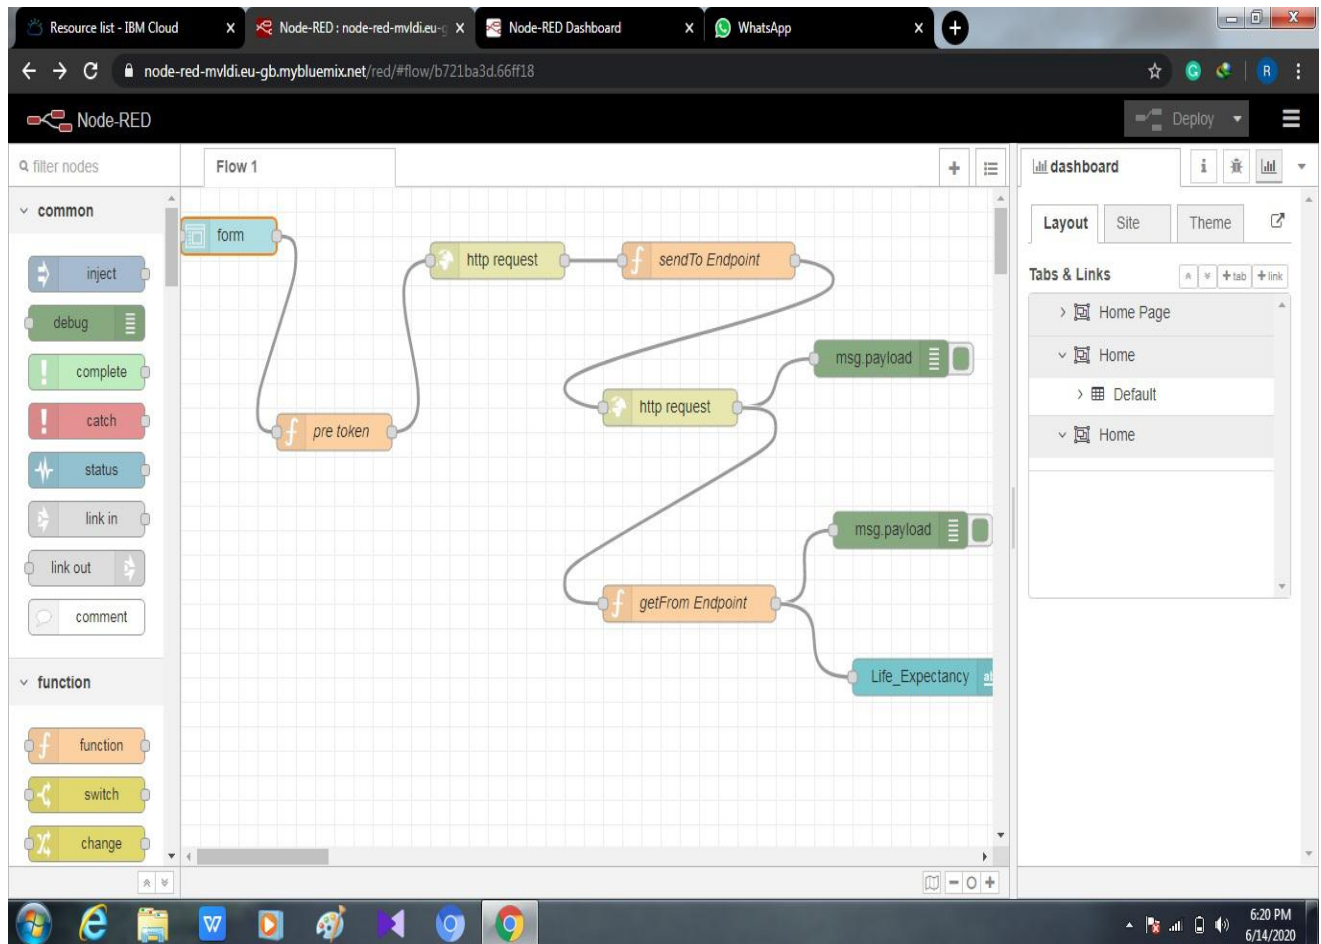
ii)



C) IBM Cloud Project Details



D) Node-Red Flow



E) Life Expectancy Prediction UI

Resource list - IBM Cloud x Node-RED : node-red-mvldi.eu- x Node-RED Dashboard x WhatsApp x

node-red-mvldi.eu-gb.mybluemix.net/ui/#/0?socketid=dmeG8SIXK55_4DgDAAAx

Home

Default

Life_Expectancy **64.8**

Year *
2015

Status("1"="Developed","0"="Developing") *
0

Adult Mortality *
263

infant deaths *
62

Alcohol *
0.01

percentage expenditure *
71.279624

Hepatitis B *
65

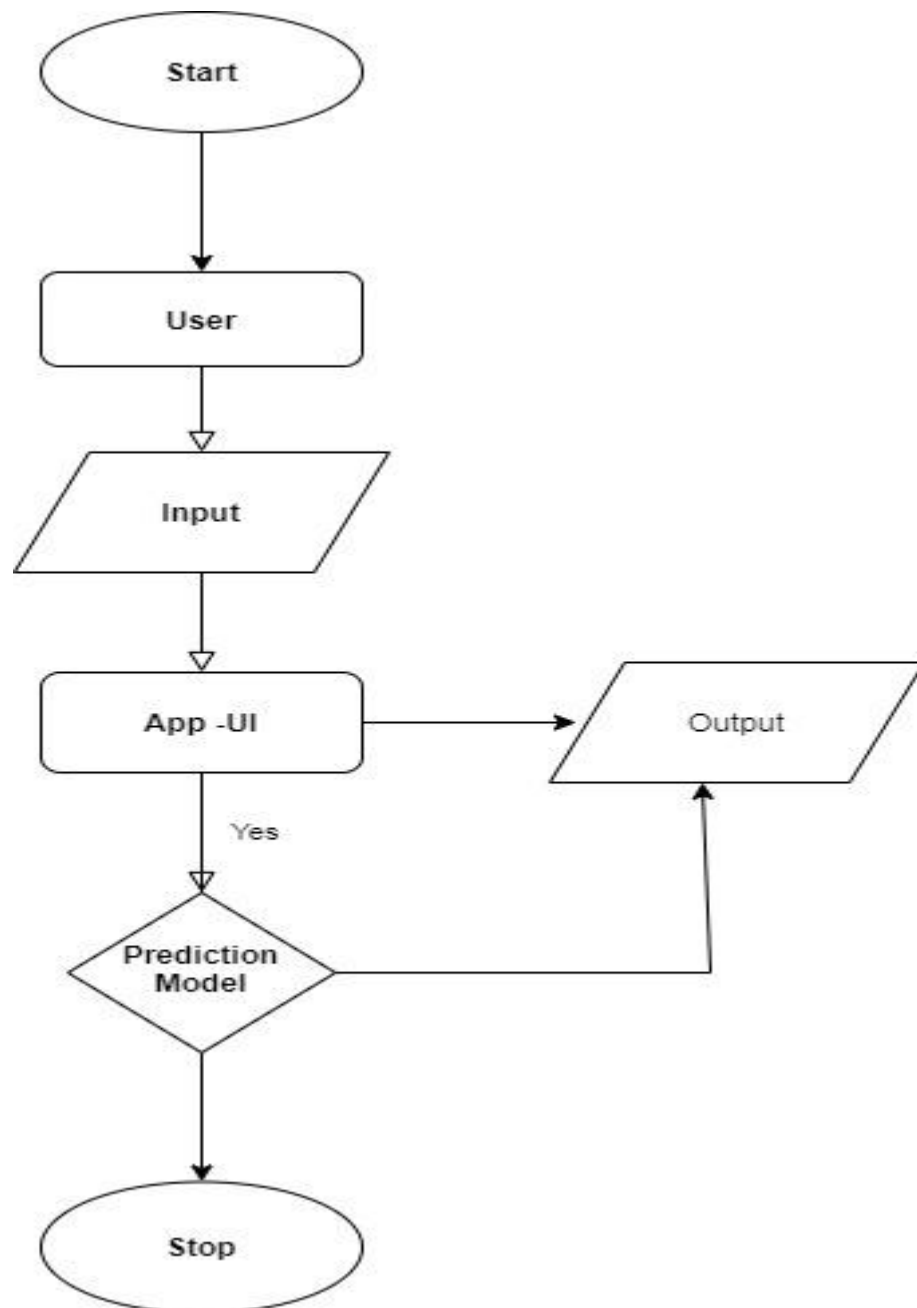
Measles *
1154

BMI *
19.1

Life_expectancy.ipynb ^ Life_expectancy.ipynb ^ Show all x

6:18 PM
6/14/2020

5. FLOWCHART



6. RESULT

The screenshot shows a web browser window with the Node-RED Dashboard open. The dashboard has a blue header with the word "Home". Below the header, there is a form titled "Default". The form contains several input fields, each with a label and a value. The values are as follows:

Field	Value
Life_Expectancy	64.8
Year *	2015
Status("1" = "Developed", "0" = "Developing") *	0
Adult Mortality *	263
infant deaths *	62
Alcohol *	0.01
percentage expenditure *	71.279624
Hepatitis B *	65
Measles *	1154
BMI *	19.1

At the bottom of the dashboard, there is a taskbar with several icons, including a folder, a document, and a web browser. The system clock in the bottom right corner shows 6:18 PM on 6/14/2020.

The screenshot shows the same Node-RED Dashboard as the previous one, but with a different set of values in the "Default" form. The values are as follows:

Field	Value
BMI	19.1
under-five deaths *	83
Polio *	6
Total expenditure *	8.16
Diphtheria *	65
HIV/AIDS *	0.1
GDP *	584.25921
Population *	33736494
thinness 1-19 years *	17.2
income composition of resources *	0.479
Schooling *	10.1

At the bottom of the form, there are two buttons: "SUBMIT" and "CANCEL". The system clock in the bottom right corner shows 6:19 PM on 6/14/2020.

7. ADVANTAGES AND DISADVANTAGES

ADVANTAGES:

- a) Health Inequalities: Life expectancy has been used nationally to monitor health inequalities of a country.
- b) Reduced Costs: This is a simple webpage and can be accessed by any citizen of a country to calculate life expectancy of their country and doesnot required any kind of payment neither for designing nor for using.
- c) User Friendly Interface: This interface requires no background knowledge of how to use it. It's a simple interface and only ask for required values and predict the output.

DISADVANTAGES:

- a) Wrong Prediction: As it depends completely on user, so if user provides some wrong values then it will predict wrong value.
- b) Average Prediction: The model predicts average or approximate value with 97.07% accuracy but not accurate value.

8. APPLICATION

- a) It can be used to monitor health inequalities of a country.
- b) It can be used to develop statistics for country development process.
- c) It can be used to analyse the factors for high life expectancy.
- d) It is user friendly and can be used by anyone.

9. CONCLUSION

This user interface will be useful for the user to predict life expectancy value of their own country or any other country based on some required details such as GDP, BMI, Year, Alcohol Intake, Total expenditure and etc.

10. FUTURE SCOPE

Future Scope of the Model can be:

a) Feature Reduction

It requires much more data about 21 columns to be known prior for predicting life expectancy which can be again difficult for a normal user to gather such data so I have decided to do some kind of feature reduction or replacement of some features as individuals or groups to make it more user friendly.

b) Attractive UI

It is a simple webpage only asking inputs and predict output. In future I have decided to make it more user friendly by providing some useful information about the country in the webpage itself so that user does not need to do any kind of prior research for the values.

c) Integrating with services such as speech recognition

11. BIBLIOGRAPHY

- <https://cloud.ibm.com/docs/overview?topic=overview-what-is-platform>
- <https://developer.ibm.com/tutorials/how-to-create-a-node-red-starter-application>
- <https://nodered.org>
- <https://www.youtube.com/embed/r7E1TJ1HtM0>
- <https://www.kaggle.com/kumarajarshi/life-expectancy-who>
- <https://bookdown.org/caoying4work/watsonstudio-workshop/jn.html>

APPENDIX

A. Source code

Importing Neccessary packages

```
In [1]: import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sn  
import missingno
```

Importing Data set

```

In [2]: import types
import pandas as pd
from botocore.client import Config
import ibm_boto3

def __iter__(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
client_c54edf641c5244a980cb66d9b7dabff7 = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='gxTpRMtR5VbP9zPeeMLTaieydKR9gKUnNo5hd0xNCPQi',
    ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3.eu-geo.objectstorage.service.networklayer.com')

body = client_c54edf641c5244a980cb66d9b7dabff7.get_object(Bucket='lifeexpectancy-donotdelete-pr-thiqdsdvtqyg3y',Key='Life.csv')['Body']
# add missing __iter__ method, so pandas accepts body as file-like object
if not hasattr(body, "__iter__"): body.__iter__ = types.MethodType( __iter__,
body )


df = pd.read_csv(body)
df.head()

```

Out[2]:

	Country	Year	Status	Life expectancy	Adult Mortality	infant deaths	Alcohol	percentage expenditure	Hepatitis B
0	Afghanistan	2015	Developing	65.0	263.0	62	0.01	71.279624	65.0
1	Afghanistan	2014	Developing	59.9	271.0	64	0.01	73.523582	62.0
2	Afghanistan	2013	Developing	59.9	268.0	66	0.01	73.219243	64.0
3	Afghanistan	2012	Developing	59.5	272.0	69	0.01	78.184215	67.0
4	Afghanistan	2011	Developing	59.2	275.0	71	0.01	7.097109	68.0

5 rows × 22 columns



In [3]: df.columns

```

Out[3]: Index(['Country', 'Year', 'Status', 'Life expectancy ', 'Adult Mortality',
    'infant deaths', 'Alcohol', 'percentage expenditure', 'Hepatitis B',
    'Measles ', ' BMI ', 'under-five deaths ', 'Polio', 'Total expenditure',
    'Diphtheria ', ' HIV/AIDS', 'GDP', 'Population',
    ' thinness 1-19 years', ' thinness 5-9 years',
    'Income composition of resources', 'Schooling'],
    dtype='object')

```

Renaming some column name

```
In [4]: df.rename(columns = {'Life expectancy ': 'Life_expectancy', 'thinness 5-9 years': 'thinness_5-9_years'}, inplace=True)
df.columns
```


```
Out[4]: Index(['Country', 'Year', 'Status', 'Life_expectancy', 'Adult Mortality',
              'infant deaths', 'Alcohol', 'percentage expenditure', 'Hepatitis B',
              'Measles ', ' BMI ', 'under-five deaths ', 'Polio', 'Total expenditure',
              'Diphtheria ', ' HIV/AIDS', 'GDP', 'Population',
              'thinness 1-19 years', 'thinness_5-9_years',
              'Income composition of resources', 'Schooling'],
              dtype='object')
```

```
In [5]: df.head()
```

Out[5]:

	Country	Year	Status	Life_expectancy	Adult Mortality	infant deaths	Alcohol	percentage expenditure	Hepat
0	Afghanistan	2015	Developing	65.0	263.0	62	0.01	71.279624	6
1	Afghanistan	2014	Developing	59.9	271.0	64	0.01	73.523582	6
2	Afghanistan	2013	Developing	59.9	268.0	66	0.01	73.219243	6
3	Afghanistan	2012	Developing	59.5	272.0	69	0.01	78.184215	6
4	Afghanistan	2011	Developing	59.2	275.0	71	0.01	7.097109	6


5 rows × 22 columns



```
In [6]: df.describe()
```

Out[6]:

	Year	Life_expectancy	Adult Mortality	infant deaths	Alcohol	percentage expenditure	Hepat
count	2938.000000	2928.000000	2928.000000	2938.000000	2744.000000	2938.000000	2385
mean	2007.518720	69.224932	164.796448	30.303948	4.602861	738.251295	80
std	4.613841	9.523867	124.292079	117.926501	4.052413	1987.914858	25
min	2000.000000	36.300000	1.000000	0.000000	0.010000	0.000000	1
25%	2004.000000	63.100000	74.000000	0.000000	0.877500	4.685343	77
50%	2008.000000	72.100000	144.000000	3.000000	3.755000	64.912906	92
75%	2012.000000	75.700000	228.000000	22.000000	7.702500	441.534144	97
max	2015.000000	89.000000	723.000000	1800.000000	17.870000	19479.911610	99



checking data types of all columns


In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2938 entries, 0 to 2937
Data columns (total 22 columns):
Country                2938 non-null object
Year                  2938 non-null int64
Status                2938 non-null object
Life_expectancy       2928 non-null float64
Adult Mortality       2928 non-null float64
infant deaths         2938 non-null int64
Alcohol               2744 non-null float64
percentage expenditure 2938 non-null float64
Hepatitis B          2385 non-null float64
Measles              2938 non-null int64
BMI                  2904 non-null float64
under-five deaths    2938 non-null int64
Polio                2919 non-null float64
Total expenditure    2712 non-null float64
Diphtheria           2919 non-null float64
HIV/AIDS             2938 non-null float64
GDP                  2490 non-null float64
Population           2286 non-null float64
  thinness 1-19 years  2904 non-null float64
thinness_5-9_years    2904 non-null float64
Income composition of resources 2771 non-null float64
Schooling            2775 non-null float64
dtypes: float64(16), int64(4), object(2)
memory usage: 505.0+ KB
```

Showing empty values in columns

```
In [8]: missingno.matrix(df)
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7fac72c1a128>
```



```
In [9]: df.isnull().sum()
```

```
Out[9]: Country          0
Year          0
Status         0
Life_expectancy    10
Adult Mortality   10
infant deaths     0
Alcohol         194
percentage expenditure    0
Hepatitis B      553
Measles         0
BMI            34
under-five deaths    0
Polio           19
Total expenditure   226
Diphtheria        19
HIV/AIDS         0
GDP            448
Population       652
thinness 1-19 years   34
thinness_5-9_years   34
Income composition of resources  167
Schooling        163
dtype: int64
```

Filling the values in columns

```
In [10]: for col in df.columns:
          df[col]=df[col].fillna(method="bfill")
          df.shape
```

Out[10]: (2938, 22)


```
In [11]: df.isnull().sum()
```

```
Out[11]: Country      0
          Year        0
          Status      0
          Life_expectancy  0
          Adult Mortality  0
          infant deaths  0
          Alcohol      0
          percentage expenditure  0
          Hepatitis B  0
          Measles      0
          BMI          0
          under-five deaths  0
          Polio        0
          Total expenditure  0
          Diphtheria   0
          HIV/AIDS     0
          GDP          0
          Population   0
          thinness 1-19 years  0
          thinness_5-9_years  0
          Income composition of resources  0
          Schooling    0
          dtype: int64
```

Feature Selection


```
In [12]: plt.figure(figsize=(20,20))
sn.heatmap(df.corr(),annot=True)
```

Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7fac729fdeb8>



Visualization of all columns

```
In [13]: sn.pairplot(df)  
plt.show()
```



Dropping some columns

```
In [14]: df = df.drop(['Country', 'thinness_5-9_years'],axis = 1)
df.head()
```

Out[14]:

	Year	Status	Life_expectancy	Adult Mortality	infant deaths	Alcohol	percentage expenditure	Hepatitis B	Measles
0	2015	Developing	65.0	263.0	62	0.01	71.279624	65.0	1154
1	2014	Developing	59.9	271.0	64	0.01	73.523582	62.0	492
2	2013	Developing	59.9	268.0	66	0.01	73.219243	64.0	430
3	2012	Developing	59.5	272.0	69	0.01	78.184215	67.0	2787
4	2011	Developing	59.2	275.0	71	0.01	7.097109	68.0	3013

Converting string values to integer(Binary-0,1)

```
In [16]: df['Status'] = (df['Status'] == 'Developed').astype(int)
df['Status'].head()
```

Out[16]:

```
0    0
1    0
2    0
3    0
4    0
Name: Status, dtype: int64
```

```
In [18]: x=df.drop(['Life_expectancy'],axis=1)
y=df.Life_expectancy
```

```
In [19]: x.head()
```

Out[19]:

	Year	Status	Adult Mortality	infant deaths	Alcohol	percentage expenditure	Hepatitis B	Measles	BMI	under-five deaths	Polio
0	2015	0	263.0	62	0.01	71.279624	65.0	1154	19.1	83	6.0
1	2014	0	271.0	64	0.01	73.523582	62.0	492	18.6	86	58.0
2	2013	0	268.0	66	0.01	73.219243	64.0	430	18.1	89	62.0
3	2012	0	272.0	69	0.01	78.184215	67.0	2787	17.6	93	67.0
4	2011	0	275.0	71	0.01	7.097109	68.0	3013	17.2	97	68.0

Importing and Training of Model

```
In [20]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.4)
```

```
In [21]: from sklearn.ensemble import RandomForestRegressor
model=RandomForestRegressor()
model.fit(x_train,y_train)
```

/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/ensemble/forest.py:246: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

```
Out[21]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=None,
                                oob_score=False, random_state=None, verbose=0, warm_start=False)
```

Testing for inputs

```
In [22]: x_test.shape
```

```
Out[22]: (1176, 19)
```

```
In [52]: y_pred=model.predict(x_test)
print(y_pred)
```

```
[69.98 68.29 84.78 ... 77.18 71.21 71.98]
```

Accuracy of Model

```
In [24]: model.score(x_train,y_train)
```

```
Out[24]: 0.990564037740675
```

```
In [25]: model.score(x_test,y_test)
```

```
Out[25]: 0.9486506073424641
```

Predicting Result

```
In [53]: Result = model.predict([[2015,0,263.0,62,0.01,71.279624,65.0,1154,19.1,83,6.0,
8.16,65.0,0.1,584.259210,33736494.0,17.2,0.479,10.1]])
print("Life Expectancy of this Country is : ",Result)
```

```
Life Expectancy of this Country is : [64.8]
```

Code for Generating Scoring Endpoint Url

```
In [27]: from watson_machine_learning_client import WatsonMachineLearningAPIClient
```

```
2020-06-14 09:52:45,012 - watson_machine_learning_client.metanames - WARNING
- 'AUTHOR_EMAIL' meta prop is deprecated. It will be ignored.
2020-06-14 09:52:48,683 - watson_machine_learning_client.wml_client_error - W
ARNING - Deployment creation failed. Error: 402. {"trace":"19fq370oimcz","err
ors":[{"code":"deployments_plan_limit_reached","message":"Current plan 'lite'
only allows 5 deployments"}]}
```

```
In [28]: wml_credentials={
    "apikey": "oA2xpCtFiQPRpgoeYUGpr_fD6ck_tdl7a1w-vD7MNj9D",
    "instance_id": "ee299d7d-5851-45a4-bf9e-87063cce7323",
    "url": "https://eu-gb.ml.cloud.ibm.com",
    "username": "18b33122-f157-4bb2-bb23-7aca666aef8e"
}
```

```
In [29]: client = WatsonMachineLearningAPIClient( wml_credentials )
```

```
In [41]: model_props = {client.repository.ModelMetaNames.AUTHOR_NAME: "rohit",
    client.repository.ModelMetaNames.AUTHOR_EMAIL: "rtiwari432@gmai
l.com",
    client.repository.ModelMetaNames.NAME: "Life_Expectancy"}
```

```
In [42]: model_artifact = client.repository.store_model(model,meta_props = model_props)
```

```
In [43]: published_model_uid = client.repository.get_model_uid(model_artifact)
```

```
In [44]: published_model_uid
```

```
Out[44]: 'e77eb785-a419-483a-b320-99cd17885abb'
```



```
In [45]: deployment = client.deployments.create(published_model_uid,name = "Life_expectancy")
```

```
#####
#####
```

Synchronous deployment creation for uid: 'e77eb785-a419-483a-b320-99cd17885abb' started

```
#####
#####
```

```
INITIALIZING
DEPLOY_SUCCESS
```

```
-----
-----
Successfully finished deployment creation, deployment_uid='5fc4b71a-21c4-4d22-ac07-82bd9d71178c'
-----
-----
```

```
In [47]: scoring_endpoint = client.deployments.get_scoring_url(deployment)
```

```
In [48]: scoring_endpoint
```

```
Out[48]: 'https://eu-gb.ml.cloud.ibm.com/v3/wml_instances/ee299d7d-5851-45a4-bf9e-87063cce7323/deployments/5fc4b71a-21c4-4d22-ac07-82bd9d71178c/online'
```

```
In [ ]:
```